

Leveraging the ‘Black Box’: Random Forest Adjustment for Causal Inference

Miles D. Williams *University of Illinois at Urbana-Champaign*

Introduction

Over the past several years, a literature centered on the application of machine learning techniques to causal identification has grown rapidly. These efforts have generated promising new approaches for those eager for more flexible and robust techniques for recovering causal estimates with observational data. Among the myriad methods that have been considered, Random Forest Regression (RF) has become increasingly popular relative to alternatives. Since its development by Breiman (2001), RF has become one of the most widely used machine learning predictors in use across a range of scientific fields—due in no small part to its high level of predictive performance and simultaneous resistance to overfitting (Wager 2016).

However, RF to-date has been of limited utility for political scientists. Its predictive performance notwithstanding, RF is a “black-box,” scoring low on interpretability—a real handicap in a discipline that prizes interpretable marginal effects and test statistics. For this reason, RF is rarely applied by political scientists, and what applications do appear in the literature are restricted to exercises in prediction or measurement. Hill and Jones (2014), for instance, use RF to test the predictive power of various political, economic, and social conditions for state repression. Bonica (2018) uses RF to assess the predictive power of campaign contributions on roll call votes in the U.S. Congress. Meanwhile, Carroll and Kenkel (2019) use RF, in addition to other machine learning algorithms, to generate a proxy for relative military power among states.

It would be a misconception, however, to suppose that RF is unusable for identifying causal relationships. For example, given the nonparametric approach taken by RF, it comes as little surprise that it has found a home among those who study the estimation of heterogeneous treatment effects (e.g., Athey and Imbens 2015; Foster et al. 2011; Su et al. 2009; Zeileis et al. 2008). These studies are joined by a broader literature that has developed creative applications of black-box machine learning techniques to causal inference (for examples see Wager 2016).

These new applications aside, it seems unlikely that political science will trade its workhorse

methods, like linear and generalized linear regression models and matching, for approaches like RF any time soon. Not all research questions demand estimation of heterogeneous treatment effects, and the benefit of learning new machine learning techniques relative to the costs in time and energy required to become familiar with them is often perceived as low.

In the face of such hesitancy, I develop a novel application of RF that leverages the power of its black-box (nonparametric) characteristics, while permitting estimation of meaningful (and interpretable) marginal effects. The approach I lay out here, which I call Random Forest Adjustment (RFA), leverages the simultaneous predictive power, and resistance to overfitting, of RF to preprocess observational data prior to estimating an average treatment effect (ATE). The approach proceeds in three simple steps: (1) compute and collect the residuals of out-of-bag (OOB) predictions of a RF model fitted using the outcome of interest as the response and a set of confounding covariates as predictors; (2) repeat setp 1, but with the causal variable of interest as the response and the confounding covariates as predictors; (3) regress the residual variation in the outcome variable on that of the causal variable and use the estimated coefficient on the treatment variable residuals as the ATE.

As I will demonstrate, this approach offers a robust (unbiased and efficient) estimate of the ATE as compared with two well-validated alternative (and more common) approaches: optimal full matching (using Mahalanobis Distance), and an interactive multiple regression model, as proposed by Lin (2013).

However, as I further demonstrate, the improved accuracy in estimated ATEs generated by RFA comes at the cost of a marginal decline in generalizability. As Aronow and Samii (2016) highlight, studies that rely on a relatively comprehensive nominal sample in observational settings generate causal effects using only a limited effective sample when accounting for regression weights. OLS, for instance, gives greater weight to observations not well explained by other covariates included as predictors in a linear model.

We should not bemoan this feature of RFA and other adjustment strategies. Having an effective sample composed of comparable observations is necessary for estimating internally valide causal estimates. The improved accuracy of RFA necessarily comes at the cost of a slightly more limited effective sample—a more comprehensive effective sample would in fact suggest that RFA poorly adjusts for variation in the data accounted for by confounding covariates.

Even so, limited generalizability notwithstanding, analysis demonstrates that the diminished overlap between nominal and effective samples produced by RFA is not overwhelming.

In the following section I offer an overview of the approach. I then validate it via a series of Monte Carlo simulations, comparing its performance to other other approaches shown be effective: optimal full matching, and the interactive multiple regression model specification suggested by Lin (2013). The results highlight the robustness of RFA to *complex confounding*—nonlinear and complex functional forms in the relationship between confounding covariates and a treatment and response. Further, the analysis reveals greater generalizability of RFA estimates of treatment effects compared to alternative approaches. As Aronow and Samii (2016) argue, when researchers use multiple regression to recover causal estimates with observational data, the effective sample produced by regression weights is often limited relative to the nominal sample—the full sample used to estimate the linear regression model. While RFA similarly relies on a limited effective sample to recover marginal effects, the RFA effective sample has greater overlap with the nominal sample than either full matching or multiple regression. RFA estimates, while still being *local*, generalize to a broader set of the nominal sample than with other approaches.

I end with an application RFA to a replication of Nielsen et al. (2011) who assess the causal effect of foreign aid shocks on the likelihood of violent armed conflict. RFA estimates differ in important ways from those identified by the authors. Nielsen et al. (2011) find that negative shocks increase the likelihood of conflict, while positive shocks have no effect. Conversely, while RFA estimates also indicate a positive effect of negative shocks on the likelihood of conflict, they identify a significant and negative estimate for the effect of positive shocks on conflict. The results obtained via RFA square better with existing theory on the role of power parity on likelihood of war (see Reed 2003). In addition, assessment of the degree of overlap between the effective and nominal samples associated with RFA and alternative estimators supports the greater generalizability of RFA estimates.

Random Forest Adjustment

RFA, like other approaches to adjusting for the influence of confounding covariates, is applicable in settings where one wishes to estimate the ATE of some intervention or causal variable, but

random assignment is in doubt. That is, one suspects propensity to receive some “treatment” to be dependent on unit characteristics of individual observations. Such is often the case in observational settings where treatment was not randomly assigned by a researcher, or else in field experiments where infelicities in design put random assignment in doubt.

Various strategies for dealing with this problem exist. Matching and its many permutations are one example, while regression-based approaches comprise another. However, these approaches, as described above, usually provide only local average treatment effects (LATEs). As a result, while being internally valid in many settings, both matching and multiple regression raise concerns about external validity.

Parametric regression models also pose additional problems. For example, in a setting with panel data, multiple regression may fail to account for heterogeneity in the confounding influence of covariates between clusters. Further, multiple regression may fail to account for nonlinearities or interactions in the d.g.p. Without perfect knowledge about the true d.g.p. under investigation, a researcher often does not know how to specify the “correct” model. Consequently, model misspecification may introduce bias into the estimated causal effect of interest—even if the research has “controlled” for all relevant confounding variables.

RFA provides an alternative adjustment strategy for generating an ATE from observational data. The procedure consists of three steps. It begins by using the out-of-bag (OOB) prediction for a response (y_i) and some causal variable of interest (z_i) as a function of covariates $x_{ik} \in X_i$ where $y_i, z_i \not\perp X_i$. The variables y_i and z_i denote vectors of response and treatment values for the i^{th} observation where $i \in I = \{1, 2, \dots, n\}$.

Because z_i and y_i are not independent of X_i , the estimated slope coefficient (α_1) from the following linear model estimated via OLS will not reflect the true ATE of z_i on y_i :

$$y_i = \alpha_0 + \alpha_1 z_i + \epsilon_i. \quad (1)$$

RFA’s solution to this problem is to partial out the variation in y_i and z_i explained by X_i prior to estimating the ATE as follows:

1. First, fit y_i as a function of X_i via a RF model and estimate the OOB observation-specific residuals:

$$\begin{aligned}\hat{y}_i &= \hat{f}(X_i), \\ \hat{y}_i^\varepsilon &= y_i - \hat{y}_i.\end{aligned}\tag{2}$$

2. This step is then repeated for z_i :

$$\begin{aligned}\hat{z}_i &= \hat{g}(X_i), \\ \hat{z}_i^\varepsilon &= z_i - \hat{z}_i.\end{aligned}\tag{3}$$

3. Finally, the ATE, adjusting for the confounding influence of X_i , is obtained by estimating the following linear model via OLS:

$$\hat{y}_i^\varepsilon = \beta_0 + \beta_1 \hat{z}_i^\varepsilon + \mu_i.\tag{4}$$

The estimate for β_1 in the above denotes the ATE.

One advantage of RFA is that it sidesteps incidental functional form assumptions that parametric adjustment strategies impose. This can be seen by considering the following fact about a linear regression model as estimated with OLS.

With the covariates given in the preceding section, define the matrix \mathbf{W} as

$$\mathbf{W} = \begin{bmatrix} 1 & z_1 & x_1^1 & \cdots & x_1^k \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & z_n & x_n^1 & \cdots & x_n^k \end{bmatrix}_{n, k+2}\tag{5}$$

and define the matrix \mathbf{y} as

$$\mathbf{y} = \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix}_{n, 1}.\tag{6}$$

Using OLS, we estimate the linear model

$$\mathbf{y} = \mathbf{W}\boldsymbol{\beta} + \boldsymbol{\epsilon},$$

where the solution for the parameters to be estimated is

$$\hat{\beta} = \begin{bmatrix} \hat{\beta}_0 \\ \hat{\beta}_1 \\ \hat{\beta}_2^1 \\ \vdots \\ \hat{\beta}_{k+2}^k \end{bmatrix} = (\mathbf{W}^\top \mathbf{W})^{-1} (\mathbf{W}^\top \mathbf{y}). \quad (7)$$

The parameter $\hat{\beta}_1 \in \hat{\beta}$ is the estimated effect of z_i (the causal variable of interest in our running example).

This approach is potentially problematic for the following reason. Consider an equivalent way of generating $\hat{\beta}_1$ via OLS.

1. Estimate a linear model excluding z_i from the matrix \mathbf{W} and with \mathbf{y} as the left-hand side variable. Denote the new $n \times k + 1$ matrix that excludes the causal variable interest as \mathbf{X} . After estimating this model using the same procedure outlined above, generate predicted values of the outcome variable and estimate the model residuals (note the similarity to step 1 of RFA).

$$\begin{aligned} \hat{\mathbf{y}} &= \mathbf{X}\hat{\gamma}, \\ \hat{\mathbf{y}}^\epsilon &= \mathbf{y} - \hat{\mathbf{y}}. \end{aligned} \quad (8)$$

2. Repeat the above, but have the $n \times 1$ matrix \mathbf{z} (a matrix of values for the causal variable) be the left-hand side variable (again, note the similarity to step 2 of RFA).

$$\begin{aligned} \hat{\mathbf{z}} &= \mathbf{X}\hat{\delta}, \\ \hat{\mathbf{z}}^\epsilon &= \mathbf{z} - \hat{\mathbf{z}}. \end{aligned} \quad (9)$$

3. Finally, estimate β_1 via OLS where the linear model is specified as

$$\hat{y}_i^\epsilon = \mu + \beta_1 \hat{z}_i^\epsilon + \epsilon_i. \quad (10)$$

The $\hat{\beta}_1$ generated from step 3 will be equivalent to the one estimated above.

At this point, it (hopefully) should be clear why a linear regression model is potentially problematic. This parametric approach, while imposing the assumption that \hat{y}_i^c and \hat{z}_i^c are linearly related, also imposes the assumption that the confounding variables are linearly and additively associated with y_i and z_i . And this is only the first layer of an onion of linear additive models within linear additive models which could be used to estimate equivalent parameters for all of the variables contained in \mathbf{W} . This is not to mention the googleplex of potential “sub”-parameters implied by assuming linearity and additivity *all the way down*. If in any of the linear additive models upon linear additive models is misspecified, this portends bias in $\hat{\beta}_1$.

This is one place where RFA shines. RFA imposes no functional form on the relationship between the confounders and the response and causal variable, and it also imposes no functional form assumptions about how any and all of the confounders relate to each other. It sidesteps the Russian doll altogether as it isolates the residual variation in the response and causal variable.

Of course, other approaches impose minimal functional form assumptions as well—matching on a distance metric, for example. And the potential issues with parametric regression-based adjustment notwithstanding, in many applied settings methods like multiple regression and matching have strong internal validity. However, as Aronow and Samii (2016) demonstrate, in providing internally valid causal estimates, both regression and matching often only provide *local* average treatment effects. Drawing a distinction between “nominal” and “effective” samples, Aronow and Samii (2016) point out that multiple regression differentially weights observations included in a nominal sample (the total set of cases a researcher may wish to generalize to), generating estimates for a causal variable of interest that only apply to a smaller effective sample (based on which observations receive greatest weight by multiple regression). As they show, the effective sample can differ markedly from the nominal sample, which restricts the generalizability of causal estimates. Matching similarly limits the scope of cases to which causal estimates apply by reducing the number of observations in the data to those that can be appropriately matched.

While internal validity is useful when the goal is merely to conduct some critical test of a hypothesis, researchers often want to know how well identified causal effects apply to a broader population. Here, again, is an area where RFA shines. Though RFA, like multiple regression and matching, generates estimates from an effective sample that, unfortunately, is more narrow than

the nominal sample, as I demonstrate in the following section, the effective sample RFA uses has greater overlap with the nominal sample than both regression and matching. Causal estimates still apply to a limited sample, but nevertheless generalize to a broader set of cases.

A Simulation

Design

To demonstrate the performance of RFA relative to alternative approaches, I conduct a Monte Carlo experiment using a pre-determined d.g.p. for a continuous response variable y_i for $i = 1, 2, \dots, n$ observations. The details of the analysis are as follows:

1. I set the number of observations to $n = 500$, and I set the ATE of a causal variable of interest $z_i \in \{0, 1\}$ to $\beta = 5$.
2. I define some continuous confounding variable x_i where $x_i \sim \mathcal{N}(50, \sigma = 10)$.
3. I define the probability that i is assigned to treatment, where $z_i = 1$ denotes treatment and $z_i = 0$ denotes control, as

$$\Pr(z_i = 1 | x_i) = \frac{1}{1 + \exp\{2 - 0.05(1.15x_i - \bar{x})^2\}}, \quad (11)$$

where the realized treatment assignment for i is determined by a random draw from a Bernoulli process given the probability of treatment for unit i as defined above.

4. Finally, values for some continuous outcome y_i are given as,

$$y_i = 1 + \beta z_i + 0.5x_i + x_i^2 + \mu_i : \mu_i \sim \mathcal{N}(0, \sigma = 10). \quad (12)$$

By design, assignment to treatment is not independent of unit characteristics. The value of x_i per each unit both determines i 's propensity to receive treatment and affects variation in the response y_i .

Additionally, to demonstrate the robustness of RFA to complex functional forms, propensity to receive treatment and values of the response are a nonlinear function of x_i . Figure 1 depicts the relationship between x_i and probability of treatment and the conditional mean of the response. As




Figure 1: The nonlinear confounding influence of x_i on treatment assignment and the conditional mean of the response.

the figure shows, the probability that i receives treatment is a convex curvilinear function of x_i . Further, the conditional mean of the response variable y_i is an exponential function of x_i .

Analysis

With the details of the simulated d.g.p. summarized, I now turn to the methods I apply to recover and evaluate estimates of β . To keep comparisons simple, I apply three approaches to obtaining β estimates: multiple regression-based adjustment, matching, and RFA.

For regression-based adjustment, I estimate the following linear model via OLS:

$$y_i = \alpha + \beta z_i + \gamma_1 x_i + \gamma_2 [z_i \cdot (x_i - \bar{x})] + \epsilon_i.$$

This interactive specification is recommended by Lin (2013). The $\hat{\beta}$ provided by the above represents the ATE for the average observation and serves as a more robust estimate of the ATE than simply additively controlling for covariates.

For matching, I rely on optimal full matching on x_i based on Mahalanobis distance. The benefit of full matching is that it makes use of as many observations as possible in identifying treatment and control comparisons (Hansen and Klopfer 2006), while it also is optimal in producing simi-

larly matched sets of observations (Rosenbaum 1991). To obtain an estimate of the ATE, I calculate the within matched set difference between treatment and control observations.

Finally, I calculate causal estimates with RFA, using x_i to residualize both y_i and z_i with random forest regression.

To make comparisons between these three approaches, I focus on four key metrics: (1) average bias, (2) mean squared error (MSE), (3) efficiency, and (4) overlap between nominal and effective samples. The first is simply estimated as

$$\text{Bias}_m = \sum_{k=1}^K \frac{\hat{\beta}_k^m - \beta_{\text{true}}}{K},$$

for $k = 1, \dots, K$ Monte Carlo replicates for a given method m . The MSE of the estimate, further, is estimated as

$$\text{MSE}_m = \sum_{k=1}^K \frac{(\hat{\beta}_k^m - \beta_{\text{true}})^2}{K}.$$

Efficiency is measured as the proportion of Monte Carlo iterations that the true ATE falls within the 95 percent confidence interval of the estimated ATE. Estimates of parameter variance are generated with a heteroskedasticity consistent variance-covariance matrix.

To assess the extent to which parameter estimates are generated from an effective sample that generalizes to the nominal sample, I also estimate the percent overlap between the effective sample distribution, and nominal sample distribution, of x_i . Following Aronow and Samii (2016), I recover the observation specific weights produced by each of the three measures by taking the square of the observation specific residual variation in z_i after applying each of the adjustment strategies. For multiple regression, this is equal to

$$w_i^{\text{LR}} = (z_i - \hat{z}_i)^2 \quad \text{s.t.} \quad \hat{z}_i = \hat{\eta}_0 + \hat{\eta}_1 x_i + \hat{\eta}_2 [z_i \cdot (x_i - \bar{x})].$$

For matching, where M is a matrix of matched set specific indicator variables, observation weights are given as

$$w_i^M = (z_i - \hat{z}_i)^2 \quad \text{s.t.} \quad \hat{z}_i = \mathbf{M}\hat{\mu}.$$

Table 1: Performance of Adjustment Strategies

model	Bias	MSE	Coverage
RFA	1.415	14.302	0.968
Matching	6.099	68.619	1
Multiple Regression	62.663	4,685.994	0.344

Finally, RFA weights are estimated as

$$w_i^{\text{RFA}} = (z_i - \hat{z}_i)^2 \quad \text{s.t.} \quad \hat{z}_i = \hat{f}(x_i),$$

where $\hat{f}(x_i)$ denotes the OOB of random forest predictions of the conditional mean of z_i given x_i .

With these weights, I generate two comparison distributions and calculate the percent overlap between the two. The first distribution is simply the unweighted distribution of x_i —this is the nominal distribution. I then compare this to that of the effective sample generated by each of the methods—that is, the distribution of $w_i^m x_i$. With these two distributions, I estimate the overlapping coefficient, or the proportion of observations that fall within both the nominal and effective samples. This is given as

$$\text{OVL}_m = \sum_{d=1}^D \min\{\kappa(x_i)_d, \kappa(w_i^m x_i)_d\}$$

where $\kappa(\cdot)$ is the gaussian Kernel density estimator and d denotes the d^{th} bin for $d = 1, \dots, D$ bins given a predetermined bandwidth.

Results

Performance of Causal Estimates

Results are shown for $K = 500$ simulated iterations of the d.g.p. Table 1 shows the estimated bias, MSE, and coverage for multiple regression, optimal full matching, and RFA. The results are shown from best to worst in performance based on average bias.

RFA performs the best, followed by matching and multiple regression. The RFA generated estimate of the ATE displays only a slight upward bias, on average returning a $\hat{\beta}$ 1.42 units higher than the true ATE. The MSE for RFA was further 14.3, and its 95 percent CI contained the true ATE 96.8 percent of the time.

Matching, meanwhile, returned an estimate of the ATE that was, on average, 6.1 units greater than the true ATE of 5. Further, the MSE for matching was 69.62 and coverage was 100 percent, suggesting wider than expected variance in the ATE estimates.

Finally, multiple regression returned an estimate of the ATE that was, on average, several magnitudes larger than the true ATE—62.66 units greater than the true ATE of 5. The MSE, further, was a staggering 4,685.99. In addition, coverage was only about 34.4 percent, which is far less than what we would expect for 95 percent CIs.

Figure 2 shows the distribution of ATE estimates from the simulation. The results show that RFA estimates most closely center around the true ATE, while matching comes in at a close second. Meanwhile, the distribution of estimates recovered by multiple regression barely overlap with the true ATE.

Distribution of ATEs




Figure 2: Variation in ATE estimates recovered by multiple regression, optimal full matching, and RFA. Results from 500 simulated iterations. The black vertical line denotes the value of the true ATE.

Nominal vs. Effective Samples

The foregoing results highlight the robustness of RFA. However, at best, they merely demonstrate that RFA scores high on internal validity. How well do RFA estimates generalize?

The results shown in Figure 3 offer an answer to this question. Shown are the distributions of overlapping coefficients for each of the simulated iterations. The greater the overlap, the greater

the percentage of observations that fall both within the nominal and effective samples.

Overlap in Nominal & Effective Sample Distributions




Figure 3: The percentage overlap between nominal and effective samples for each method across Monte Carlo replicates.

In settings with nonrandom treatment assignment, 100 percent overlap is rarely possible. The idea that using multiple regression or matching with a large, observational dataset ensures greater external validity is a common fallacy—albeit, an understandable one. However, the effective sample that such methods use to actual estimate a causal effect often differs substantially from the nominal sample.

This much is apparent for each of the methods considered here—even for RFA. However, if external validity is a concern, RFA may be the optimal choice. In comparing the distribution of overlapping coefficients per each simulated iteration, the effective sample used by RFA has nearly double the overlap with the nominal sample relative to matching and multiple regression. On average, the effective samples from multiple regression and matching overlap with the nominal sample 10 and 12.6 percent of the time. The effective sample from RFA, meanwhile, overlaps with the nominal sample, on average, 22.1 percent of the time. This is an appreciable and statistically significant improvement in the degree to which the effective sample generalizes to the nominal.

Application to Real-World Data

To demonstrate the application of RFA, I replicate analysis done by Nielsen et al. (2011). The authors test the effect of so-called “aid shocks” on the probability of civil war in countries that receive Official Development Assistance (ODA). An aid shock is defined as a sudden change in total aid received as a proportion of gross domestic product (GDP). The authors code a decline in aid per GDP at or below the 15th percentile for the entire sample of cases in their analysis as a *negative* aid shock, and an increase in aid per GDP at or above the 85th percentile as a *positive* aid shock.

Shocks in aid flows are theorized to influence the likelihood of armed conflict between a government and potential insurgents via a number of pathways. One is that governments may use aid to supplement side-payments to opposing factions. By “buying off” potential insurgents, states hope to reduce the likelihood of conflict.

A second pathway is the impact of aid flows on power parity between the state and insurgents. A sudden decline in ODA is presumed to improve the power of insurgents relative to the state due to loss of aid revenue. Consequently, the state is left with fewer resources, improving rebels’ perceived chances of victory in armed conflict.

The authors hypothesize that negative aid shocks increase the probability of conflict. However, they offer reasons to expect either a positive or negative effect of positive aid shocks. A negative effect follows from logic that mirrors that given above: more aid means more resources for the state to use as side payments and to combat insurgents. Conversely, windfalls in aid may incite bargaining problems between the government and opposition groups. With an increase in resources comes conflict over how those resources should be distributed, and insurgents may demand more than they previously were receiving. The authors, therefore, leave the question of the effect of positive aid shocks to empirical investigation.

In their analysis, Nielsen et al. (2011) find that negative aid shocks increase the probability of armed conflict. Meanwhile, positive aid shocks have a null effect—perhaps as a result of the cross-cutting forces highlighted above. For their analysis, the authors use both regression-based adjustment (via logistic regression) and genetic matching to generate these findings. Here, I apply RFA to this question, alongside the two approaches I used in the previous simulation—optimal

full matching and multiple regression.

The Data

With the replication dataset made available by Nielsen et al. (2011), I estimate the effect of *aid shocks* on the incidence of *civil war*. After multiple imputation with random forests, the sample includes 201 aid recipient countries from 1975 to 2007 for a total of 5,130 country-year observations.

Following the authors, I control for the following confounding variables: human rights violations, number of assassinations, number of riots, number of strikes, number of anti-government demonstrations, infant mortality, number of contiguous countries experiencing a civil war, regime type, poverty, population, ethnic and religious fractionalization, the Cold War, mountainous terrain, noncontiguousness, and year. All time-varying variables are lagged by one year.

The authors, in their main specification, code an aid shock as occurring when the change in aid per GDP is less than or equal to the 15th percentile for the entire sample. Because this threshold is arbitrary, and since the estimated effect of an aid shock may be sensitive to the cutoff used, I calculate the effect of aid shocks across a range of percentile cutoffs (the 15 to the 85 percentiles in increments of 10). For all shocks coded below the 50th percentile, I code the shock as negative (all values less than or equal to the cutoff are coded as 1). Conversely, for all shocks above the 50th percentile, I code the shock as positive (all values greater than or equal to the cutoff are coded as 1). Varying the cutoff in this way allows me to observe the validity of different aid shock thresholds and to see whether negative shocks have differential effects from positive shocks.

Figure 4 shows the distribution of values of percent change in total aid received by states per GDP. The dashed vertical line denotes the 15th percentile. Inexplicably, this value differs from that calculated by Nielsen et al. Nevertheless, since I use different percentile cutoffs in the analysis, the difference between the value identified as the cutoff point for aid shocks by the authors and my own is not of much concern.

Analysis

To recover estimates of the effect of aid shocks on civil war onset, I use the same three approaches as in the previous simulation. In addition to RFA, I use optimal full matching, as well as regression-

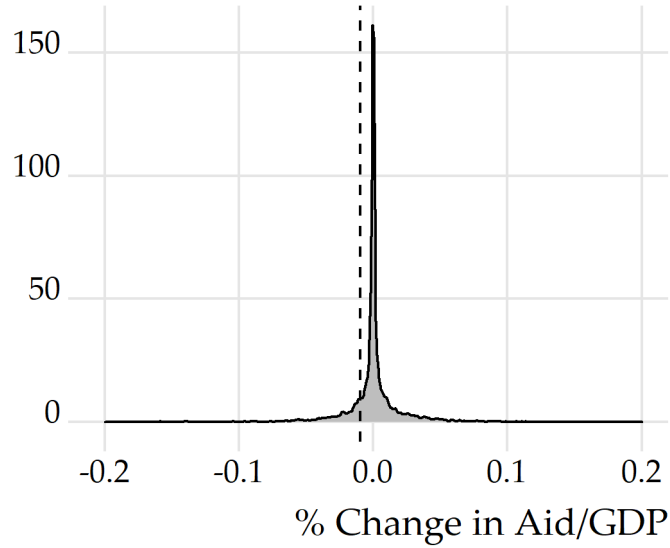


Figure 4: Distribution of the causal variable (percent change in aid/GDP). The vertical line denotes the 15th percentile.

based adjustment.¹

Figure 5 shows RFA, genetic matching, and OLS estimates of the effect of aid shocks on the probability of civil war onset. Results are shown for varying cutoff points for aid shocks. Unlike Nielsen et al. (2011), I am unable to recover evidence of a significant effect of aid shocks at the 15th percentile, regardless of method used. More puzzling, the estimate returned by matching is negative, suggesting a marginal decline in the probability of civil war. However, at the 25th percentile, both RFA and multiple regression return a significant and positive estimate of the effect of a negative shock in aid flows on the likelihood of civil war. Further, while not statistically significant, matching as well returns a positive estimate.

In terms of positive aid shocks, only RFA and matching return a significant and negative estimate for the probability of civil war, and only for increases in aid at or greater than the 75th percentile—contrary to the findings of Nielsen et al. (2011) who fail to identify a significant effect for positive aid shocks.

To assess the degree to which estimates returned by RFA, matching, and multiple regression generalize to the nominal sample of cases, I calculated the overlapping coefficient per each confounding covariate per each method. I used negative aid shocks at the 15th percentile as the

¹For matching, I use the ‘fullmatch’ function in the ‘optmatch’ package in ‘R’.




Figure 5: The effect of aid shocks on the probability of civil war. Blue estimates obtained via RFA. Red estimates obtained via optimal full matching. Green estimates obtained via multiple regression. 95 percent CIs calculated from robust standard errors. Results shown with different cutpoints for “aid shocks.”




Figure 6: The degree of overlap between the nominal and effective samples for RFA, matching, and multiple regression.

Table 2: Best Coverage across Covariates

	Best Coverage for N Covariates
Matching	3
Multiple Regression	6
RFA	12

treatment variable in calculating overlap. Estimates are shown in Figure 6. The degree of overlap between effective and nominal samples differs markedly across covariates. While some have nearly 50 percent overlap (infant mortality, the Cold War dummy, religious fractionalization, income, and political instability), one has next to zero (oil exports). Overlapping coefficients further appear correlated across estimators, suggesting that each of the methods estimates the ATE of aid shocks using roughly similar effective samples.

However, despite marginal similarities in overlap, RFA has a noticeable edge. Table 2 shows the number of covariates for which each method's overlapping coefficient is greatest. Matching only displays greatest overlap for 3 of the 21 covariates. Multiple Regression does slightly better with greatest overlap for 6 covariates. RFA, however, has the greatest overlapping coefficients for 12 of the 21 covariates—twice that of multiple regression, and four times that of matching. Though it is important to avoid making too much of these differences—imperfect and limited overlap is to be expected, not to mention necessary, for recovering estimates based on appropriate comparisons in the data—RFA nonetheless bases estimates on a slightly more representative sample, as it did in the previous simulation.

Table 3: Difference between Effective and Nominal Samples

	variable	Regression	Matching	RFA
1	HR Violations	0.11***	0.11***	0.14***
2	Assassinations	0.02***	0.02***	0.02***
3	Riots	-0.1***	-0.09***	-0.11***
4	Strikes	-0.4***	-0.37***	-0.44***
5	Demonstrations	-0.27***	-0.18***	-0.15***
6	Infant Mortality	-2.4***	-2.09***	-2.55***
7	Neighbor Civil War	0.03***	0.03***	0.04***
8	Partial Autocracy	0.05***	0.05***	0.06***
9	Partial Democracy	0.01**	0.01***	0.01***
10	Factional Democracy	-0.01.	-0.02.	-0.02**
11	Full Democracy	-0.01	0	0
12	Income (ln)	0	0.01	0.01
13	Population (ln)	-0.06***	-0.05***	-0.07***
14	Oil Exports	0.23	0.24	0.3*
15	Political Instability	0	0.02	-0.02.
16	Ethnic Fractionalization	0	0.01.	0.01
17	Religious Fractionalization	-0.03	0	-0.05**
18	Cold War	11.46***	11***	14.08***
19	Mountains	0.08***	0.09***	0.1***
20	Noncontiguous	0.05***	0.04***	0.06***
21	Year	0.03***	0.03***	0.04***